



Qumis

Strategic Messaging Document

This document provides comprehensive messaging guidance for Qumis, the coverage intelligence platform for commercial P&C insurance professionals. It is designed for immediate use by marketing, sales, customer experience and leadership teams to ensure consistent, compelling communication across all channels.

Messaging Goals

- For Qumis to be recognized as the premier provider of Coverage Intelligence
- Maintain and expand recognition of the value Qumis brings to insurance experts
- Empower Qumis employees, partners and vendors/agencies to speak with a consistent voice across all channels and audiences

1. Positioning Statement

2. Core Value Propositions

Legal-Grade Accuracy from Attorney-Trained AI

Coverage Intelligence, Not Workflow Automation

Trusted by Enterprise Leaders

Built for How Insurance Actually Works

Market Intelligence Powered by Proprietary Data

3. Messaging by Audience

Brokerage Executives

Claims Leadership

IT & Operations

4. Persona-Specific Messaging

Front Office / Producers

Back Office / Account Managers

Claims Professionals

5. Competitive Differentiation

The Competitive Landscape

vs. Workflow Automation Platforms (FurtherAI, Fulcrum)

Killer Questions to Expose the Gap:

Competitor-Specific Notes (internal use only — do not share with prospects):

vs. Generic AI (ChatGPT, etc.)

Form 1: "Why can't we just use ChatGPT?" (The casual dismissal)

Form 2: "Our IT team can build something with the API." (The build-vs-buy objection)

Form 3: "AI is improving fast — won't this be commoditized?" (The wait-and-see objection)



[vs. Outside Counsel & Coverage Consultants](#)

[vs. Status Quo \(Manual Processes\)](#)

[Rules of Engagement:](#)

[6. Proof Points & Social Proof](#)

[Enterprise Traction](#)

[Named Customer Logos](#)

[Quantified Impact Metrics](#)

[Claims-Specific Metrics \(sourced from customer interviews\)](#)

[Brokerage-Specific Metrics](#)

[Customer Testimonials](#)

[Coverage Expertise & Accuracy](#)

[Operational Efficiency & ROI](#)

[Trust & Reliability](#)

[Testimonial Selection Guide](#)

[Credibility Markers](#)

[Founding Team](#)

[Institutional Validation](#)

[Product & Technology](#)

[Case Studies \(Narrative Format\)](#)

[Case Study 1: "The \\$21,000 Exclusion" — Vance Woods, Skyward Specialty](#)

[Case Study 2: "The Capacity Transformation" — Brandon M., Top 5 Regional Broker](#)

[Case Study 3: "Land and Expand" — NFP](#)

[7. Language Guide](#)

[Key Phrases by Context](#)

[Strategic Positioning Phrases \(use in marketing, investor materials, executive conversations\)](#)

[Sales Conversation Phrases \(use in demos, discovery calls, follow-up emails\)](#)

[Competitive Phrases \(use when directly compared to alternatives\)](#)

[Customer Success Phrases \(use in onboarding, QBRs, renewal conversations\)](#)

[Competitor Reference Language](#)

[Regulatory & Legal Guardrails](#)

[Tone Principles](#)

[8. Objection Handling](#)

[Objection 1: "How is this different from ChatGPT?"](#)

[Objection 2: "We already have \[workflow automation tool\]"](#)

[Objection 3: "How do I know I can trust the AI output?"](#)

[Objection 4: "Is this going to replace my team?"](#)

[Objection 5: "What about security and compliance?"](#)

[Objection 6: "How long does implementation take?"](#)

[Objection 7: "We tried \[other AI tool\] and it didn't work"](#)



[Objection 8: "This is too expensive" / "We don't have the budget"](#)

[Objection 9: "Our IT team can build this internally"](#)

[Objection 10: "We need to see it work on our policies first"](#)

[9. The Qumis Ascend Program](#)

[Why Ascend Exists: The Evidence Base](#)

[How to Position Ascend](#)

[Ascend as Competitive Differentiator](#)

[Ascend Proof Points](#)

[Key Phrases for Ascend Conversations](#)

[10. Data, AI, and Customer Privacy: Messaging Guide](#)

[The Core Tension](#)

[The Questions Your Team Will Get](#)

[Language Discipline](#)

[Red Lines: What Your Team Must Never Say](#)

[Internal Note](#)



1. Positioning Statement

Qumis is the coverage intelligence platform for commercial P&C insurance professionals. Built by a licensed coverage attorney and AI experts, Qumis combines legal-grade policy analysis with proprietary market intelligence drawn from thousands of commercial insurance programs — delivering not just what a policy says, but what it means in the context of the broader market. Unlike generic AI tools that process documents or workflow automation platforms that move them faster, Qumis understands the legal nuance of insurance contracts and benchmarks coverage against real market data, enabling brokers, carriers, and claims teams to interpret, evaluate, and position policies with the precision of experienced coverage counsel and the market context of a seasoned industry insider — at enterprise speed and scale.

Why this positioning works:

- Category ownership: "Coverage intelligence" still anchors the category, but now encompasses both legal analysis and market intelligence — a broader, more defensible category definition that competitors can't follow you into.
- Credibility anchor: Attorney-built credentials establish domain authority on the analysis side; the proprietary data asset ("thousands of commercial insurance programs") establishes authority on the market intelligence side. Two moats, not one.
- Dual value proposition: "What a policy says" vs. "what it means in context" captures the full spectrum — from document-level analysis to market-level insight. This is the difference between a reading tool and an intelligence platform.
- Clear contrast: Now separates from three competitors simultaneously: generic AI (no domain expertise), workflow automation (no coverage reasoning), and traditional consulting/law firms (no scale, no market data). The competitive frame is wider.
- Outcome focus: "Interpret, evaluate, and position" maps to the actual jobs-to-be-done across your three user personas — claims professionals interpreting coverage, account teams evaluating adequacy, and producers positioning placements.
- Data moat signaling: Referencing proprietary market data plants the flag that this isn't just AI applied to documents — it's accumulated intelligence that gets more valuable over time. This matters for buyers (why Qumis over DIY) and for investors (why this is defensible).

2. Core Value Propositions

Legal-Grade Accuracy from Attorney-Trained AI



Qumis doesn't just read policies, it reasons about them. Our AI was developed by a licensed coverage attorney and trained on the interpretive frameworks that experienced counsel use to analyze coverage questions. Every output is citation-linked to source documents, with built-in safeguards that say "I don't know" when evidence is insufficient.

Proof Points:

- Founded by a licensed insurance coverage attorney with law firm and in-house experience
- Multi-agent AI architecture mirrors actual legal review processes
- Citation-linked outputs enable audit trails and defensible decisions
- "I don't know" responses prevent AI hallucination on ambiguous coverage questions

Messaging rationale: This is our primary differentiator. Generic AI hallucinates on legal nuance; we're built for legal-grade precision.

Coverage Intelligence, Not Workflow Automation

Competitors automate document shuffling. Qumis unlocks the intelligence inside the contract itself. We transform unstructured policy language into structured, queryable data that enables coverage reasoning at scale, turning illiquid contract intelligence into a strategic asset.

Proof Points:

- Proprietary insurance taxonomies understand clause roles (exclusions, conditions, endorsements)
- Structured data extraction enables programmatic analysis across policy portfolios
- Coverage comparison goes beyond differences to surface strategic implications
- Legal Search (beta) connects policy questions to relevant case law

Messaging rationale: Draws clear contrast with FurtherAI and other workflow tools. We're not faster document processing; we're deeper document understanding.

Trusted by Enterprise Leaders

Qumis is the platform of choice for the industry's most demanding users. Our enterprise-grade security, SOC 2 certification, and proven scalability have earned the trust of top-tier brokerages navigating their highest-stakes coverage decisions.

Proof Points:

- 5 of the top 15 U.S. insurance brokerages using Qumis
- 127% net revenue retention demonstrates expansion within accounts
- SOC 2 Type I certified with enterprise-grade data security
- NFP deployment: expanded from pilot to 250+ seats in production

Messaging rationale: Enterprise logos provide credibility that no feature list can match. This signals we're a proven, enterprise-ready solution.



Built for How Insurance Actually Works

Insurance professionals don't work in generic workflows. They navigate layered coverage programs, endorsement towers, and nuanced policy language. Qumis was purpose-built for these realities, not adapted from generic document processing tools.

Proof Points:

- Document Vault centralizes policy portfolios for cross-document analysis
- Handles complex scenarios: endorsements, amendments, layered programs
- Prompt Library enables team-wide consistency on coverage workflows
- QumisScore provides proprietary accuracy metrics for comparison results

Messaging rationale: Signals deep domain expertise that generic tools can't match.

Market Intelligence Powered by Proprietary Data

Qumis doesn't just analyze policies in isolation. Our platform draws on a proprietary database of thousands of commercial insurance programs to contextualize every analysis — surfacing how coverage compares to market norms, where gaps exist relative to industry benchmarks, and what best-in-class coverage looks like for a given risk class. This turns every policy review from a standalone reading exercise into a market-informed coverage evaluation.

Proof Points:

- Proprietary database of 10,000+ commercial insurance policies across major lines of business
- QumisScore benchmarks coverage quality against market standards
- Coverage gap detection informed by cross-portfolio pattern analysis
- Market context enriches every report, comparison, and coverage opinion

3. Messaging by Audience

This section provides messaging frameworks for organizational decision-makers and budget holders. For day-to-day practitioner messaging that drives adoption and internal advocacy, see Section 4.

Brokerage Executives

Primary concerns: Competitive differentiation, talent leverage, E&O risk reduction, client retention, organic growth from existing accounts



Lead message: "Qumis gives your team a coverage advantage your competitors can't match — legal-grade insight and market intelligence that wins accounts, retains clients, and scales your best people's expertise across the organization."

Key talking points:

- **Talent multiplier:** Senior coverage expertise scales across junior team members — every account manager and producer operates with the analytical depth of your most experienced people.
- **Market intelligence advantage:** Your team doesn't just analyze policies — they benchmark coverage against market standards across thousands of commercial programs, giving clients data-backed confidence that their program is competitive.
- **Win rate impact:** Data-driven, citation-linked coverage proposals backed by market context beat incumbent guesswork. Your producers walk into meetings with insights the competition can't access.
- **E&O protection:** Citation-linked analysis creates defensible documentation. Every recommendation traces back to source policy language and market benchmarks.
- **Client stickiness and organic growth:** Instant, authoritative coverage answers strengthen relationships. Proactive identification of coverage gaps relative to market norms creates natural upsell and cross-sell conversations that drive organic revenue growth.

Proof point to lead with: "The brokers at NFP reduced complex policy review time by 80% — and expanded from 50 to 250+ active users within months."

Messaging rationale: Executives care about competitive positioning, revenue growth, and risk reduction — in that order. The revised lead message avoids dismissing workflow efficiency (which executives are actively buying) and instead positions Qumis as delivering something competitors can't access: legal-grade analysis combined with proprietary market intelligence. The addition of organic growth and cross-sell language connects directly to the revenue expansion story that justifies platform investment.

Claims Leadership

Primary concerns: Cycle time, indemnity leakage, consistency of coverage determinations, bad-faith exposure, outside counsel spend

Lead message: "Qumis gives your claims organization the coverage reasoning capability of experienced coverage counsel — at the speed and consistency your operation demands. Every adjuster makes better coverage determinations, backed by citation-linked analysis and market intelligence that strengthens your positions."

Key talking points:



- **Faster, more accurate determinations:** Coverage analysis that took days now takes hours, with citation-linked reasoning that catches exclusions, endorsements, and coverage defenses that manual review misses under deadline pressure.
- **Consistency at scale:** Different adjusters interpreting the same policy language differently creates bad-faith exposure and regulatory risk. Qumis standardizes coverage reasoning across your entire claims organization.
- **Reduced outside counsel spend:** Coverage questions that currently get sent to outside counsel at \$500–\$800/hour can be resolved internally with attorney-grade analysis, reserving outside counsel for genuinely complex disputes.
- **Market context for coverage positions:** Contextualize coverage determinations against how similar provisions are structured across thousands of commercial programs — strengthening reservation-of-rights letters and coverage positions with market-informed analysis.
- **Defensible documentation:** Citation-backed coverage determinations create audit trails that withstand scrutiny from insureds, regulators, and bad-faith plaintiffs.

Proof point to lead with: "Built by a licensed insurance coverage attorney who understands exactly how coverage determinations get challenged. Trusted by claims teams at top-tier brokerages and carriers handling complex commercial programs."

Messaging rationale: Claims leadership buys on accuracy, consistency, and defensibility — not speed alone. The outside counsel spend reduction is a concrete, measurable ROI that finance teams understand. Market intelligence lands here as a credibility enhancer for coverage positions, which is how claims leadership actually thinks about data.

IT & Operations

Primary concerns: Security posture, integration complexity, vendor stability, implementation timeline, change management burden

Lead message: "Qumis is enterprise-ready from day one — SOC 2 certified, designed for fast deployment, and built to work within your existing technology ecosystem without requiring your team to change how they work."

Key talking points:

- **SOC 2 Type I certified:** Enterprise-grade security and compliance, purpose-built for organizations handling sensitive insurance policy and claims data.
- **Data isolation and privacy:** Client data is private and never used for AI model training. Meets the data handling standards your most demanding clients require.



- **Integration with existing ecosystems:** Designed to work alongside your existing AMS (Applied Epic, Sagitta), document management systems (ImageRight, SharePoint, Box), and email infrastructure (Outlook) — not replace them.
- **Minimal IT lift:** Cloud-based deployment with no on-premise infrastructure required. Implementation measured in weeks, not months.
- **API-first architecture:** Connects to existing systems and workflows programmatically, enabling custom integrations without disrupting established processes.

Proof point to lead with: "Trusted by 5 of the top 15 U.S. brokerages, including enterprise-wide deployments at organizations with complex security and compliance requirements."

Messaging rationale: IT buyers at brokerages have been burned by tools that require ripping out existing infrastructure or creating parallel workflows. The integration language is specific to insurance technology ecosystems (AMS, document management, Outlook) rather than generic SaaS boilerplate — signaling that Qumis understands the actual technology landscape these teams manage. Data isolation language addresses the specific concern insurance IT teams raise about AI tools using their policy data.

4. Persona-Specific Messaging

This section focuses on the day-to-day practitioners who will use Qumis. Each persona has distinct pain points, priorities, and value drivers that require tailored messaging to drive adoption and internal advocacy. Where Section 3 addresses why decision-makers should buy Qumis, this section addresses why practitioners will champion and use it.

Front Office / Producers

Primary Value Proposition

Qumis gives you a coverage advantage your competitors don't have. In minutes, you'll deliver legal-grade coverage analysis backed by market intelligence from thousands of commercial programs — presenting data-driven insights that differentiate your proposals and win accounts. While your competitors walk in with generic comparisons, you'll show prospects exactly how their coverage stacks up against the market and where the gaps are, with citations to prove it.

Key Benefits

- **Market-informed proposals:** Show prospects how their coverage compares not just to their expiring, but to what you're seeing across comparable programs in their industry. That's a fundamentally different conversation than "I reviewed your policy."



- **Competitive differentiation:** Present detailed coverage comparisons that expose gaps in incumbent programs with market context, demonstrating value prospects can't get from any other broker.
- **Revenue acceleration:** Close more deals by answering complex coverage questions instantly with citation-linked analysis, keeping prospects engaged instead of waiting days for your team to get back to them.
- **Client confidence from day one:** Citation-linked analysis backed by market benchmarks shows you've done deeper homework than anyone else. You're not making claims — you're showing data.
- **More selling, less researching:** Focus time on building relationships and closing deals instead of wading through policy documents or waiting on coverage teams to produce analysis.

Pain Points We Solve

- **Inability to demonstrate differentiated value quickly:** Without market data and legal-grade analysis to back up your recommendations, you're just another broker making claims. Prospects need proof, not promises — and your experienced competitors have deep institutional knowledge you need to match.
- **Slow sales cycles:** Manual policy reviews create bottlenecks that kill momentum in competitive situations. Prospects go cold waiting for coverage analysis that your coverage team is producing manually.
- **Lack of competitive intelligence:** You can't expose gaps in the incumbent program if you can't quickly analyze their current coverage against market standards. Generic comparisons don't win business; market-informed ones do.
- **Real-time information gap:** In client meetings, you need instant answers to coverage questions. Saying "I'll get back to you" erodes confidence and gives competitors an opening.

Proof Points to Emphasize

- **80% faster policy reviews:** NFP producers reduced complex policy review time by 80%, freeing time for client-facing activity
- **Market benchmarking:** Coverage comparisons contextualized against thousands of commercial programs across major lines of business
- **Citation-backed proposals:** Every recommendation links directly to policy language — show prospects exactly where coverage gaps exist
- **Built by attorneys:** Developed by a licensed coverage attorney who understands exactly what it takes to evaluate coverage and win business on expertise

Suggested Talk Track Opening

When Competing Against Workflow Platforms:



"Those platforms are great for speeding up proposals and data entry — that's table stakes. But when you're sitting across from a prospect and they ask 'How does my coverage compare to what you're seeing in the market?' you need real coverage intelligence, not just faster document processing. Qumis gives you the answer that wins the business — legal-grade analysis backed by market data from thousands of commercial programs, with citations you can show them on the spot. That's the competitive advantage that matters."

Cold Open for Producer Conversations:

"I know your time is valuable, and you need to walk into prospect meetings with something your competitors can't match. Qumis gives you instant, legal-grade analysis on any policy — I can show you coverage gaps and market benchmarking backed by actual policy language, right here in our meeting. You'll show prospects how their program stacks up against comparable accounts in their industry. Let me pull up a program and show you what I mean."

Back Office / Account Managers

Primary Value Proposition

Qumis makes you the most knowledgeable person on your client's coverage program. Instead of spending your days extracting data from policy documents, you'll have structured intelligence at your fingertips — with the market context to proactively identify coverage gaps, flag renewal risks, and provide the kind of strategic guidance that makes clients see you as an indispensable partner, not a service coordinator.

Key Benefits

- **Proactive client advisor:** Surface coverage gaps and market insights at renewal — "Based on what we're seeing across comparable programs, here are three areas where your client's coverage has fallen behind market standard." That's strategic account management, not just policy processing.
- **Confidence on every client call:** Citation-linked analysis means you never have to guess or hedge on a coverage question. When a client calls about a specific exclusion, you have the answer — backed by source policy language.
- **Operational efficiency:** Slash the time spent on policy reviews, coverage summaries, and renewal preparation from hours to minutes per task. Structured data extraction eliminates manual copy-paste and the typos that come with it.
- **Error reduction:** Citation-linked outputs eliminate the risk of missing critical exclusions or misreading endorsements — no more corrections to clients, no more E&O anxiety.
- **Handle renewals without the panic:** Manage growing books of business without working weekends or the 2 AM realization that you missed something in a 200-page program.



Pain Points We Solve

- **Reactive instead of strategic:** You spend your days responding to urgent requests instead of proactively strengthening client relationships. You know you should be surfacing renewal insights and coverage recommendations, but there's never enough time.
- **Manual data extraction burden:** Hours spent pulling policy details, limits, deductibles, and endorsements from PDFs into summaries and spreadsheets — tedious work that's also the most error-prone part of your job.
- **No market context for renewals:** You can tell a client what their policy says, but you can't tell them how it compares to what similar organizations are buying. That market insight is what separates strategic advisors from service teams.
- **Coverage questions that stump you:** When a client asks about a specific endorsement interaction or whether an exclusion applies to their situation, you either guess, escalate, or say "I'll get back to you." All three erode confidence.
- **Firefighting vs. strategic work:** Your days are consumed by the urgent instead of the important. Every hour you spend on manual policy extraction is an hour you're not spending on the client relationship work that actually drives retention.

Proof Points to Emphasize

- **Structured data extraction:** Qumis automatically pulls policy details into organized, searchable formats — no more copy-paste or typos
- **Market intelligence:** Benchmark client coverage against comparable programs to proactively identify gaps and strengthen renewal conversations
- **Document Vault:** Centralized repository for all client policies — instant access for the entire team, no more hunting through shared drives
- **Prompt Library:** Team-wide consistency on coverage workflows — everyone delivers the same quality, reducing training burden
- **User testimonial:** "It made our policy reviews faster, sharper, and more accurate while eliminating the need for an additional hire — it's a no-brainer." — Senior account manager at a top-15 brokerage

Integration Note for AMs:

Qumis is designed to work within the tools you already use — integrating with your email workflow through Outlook and connecting to document management systems like SharePoint and Box. You don't need to change how you work; Qumis fits into your existing process.

Suggested Talk Track Opening

When Competing Against Workflow Platforms:



"Workflow platforms are excellent for proposal generation and AMS data entry — they save you hours on admin work. But when a client calls asking about a specific exclusion, or why their claim might not be covered, or how their coverage compares to the market, you need more than automated data extraction. You need to actually understand the policy and its context. Qumis gives you the coverage intelligence to answer confidently, with citations you can share and market benchmarks to back you up."

Cold Open for AM Conversations:

"I know you're juggling dozens of accounts and every policy request feels urgent. What if you could walk into your next renewal meeting not just with the policy analysis, but with market insights showing how your client's coverage stacks up against comparable programs? Qumis handles the heavy lifting — extracting data, identifying gaps, benchmarking against the market — so you can focus on the strategic conversations that keep clients loyal. Let me show you what a renewal prep looks like with Qumis."

Claims Professionals

Primary Value Proposition

Qumis accelerates claims resolution while strengthening your coverage position. When a claim arrives, you'll instantly understand the applicable coverage, identify relevant exclusions and endorsements, and build defensible positions backed by citation-linked analysis and market intelligence. Instead of days spent manually reviewing policies and researching coverage questions, you'll deliver faster, more confident determinations with documentation that holds up under scrutiny.

Key Benefits

- **Faster cycle times:** Coverage analysis that took days now takes hours — accelerating claims resolution and getting to the right answer before deadline pressure forces a rushed decision.
- **Improved accuracy:** Citation-linked analysis ensures you never miss a critical exclusion or misinterpret an endorsement during coverage determination. Every relevant provision is surfaced, not just the ones you happened to find.
- **Market-informed coverage positions:** Contextualize your coverage determinations against how similar provisions are structured across thousands of commercial programs — strengthening reservation-of-rights letters and coverage positions with market data, not just policy language.
- **Defensible documentation:** Citation-backed coverage determinations create audit trails that withstand scrutiny from insureds, regulators, and bad-faith plaintiffs. Your reasoning is documented, sourced, and replicable.



- **Identify coverage defenses comprehensively:** Surface every applicable exclusion, limitation, and condition — including the ones that get overlooked under deadline pressure — to build the strongest defensible position on every file.

Pain Points We Solve

- **Manual policy review bottlenecks:** Each claim requires hours of reading through policies, endorsements, and exclusions before you can make coverage determinations — and the more complex the program, the more likely something gets missed.
- **Inconsistent handling across the team:** Different adjusters interpret the same policy language differently, creating exposure to bad-faith claims and regulatory scrutiny. There's no standardized approach to coverage reasoning.
- **Missed coverage defenses under pressure:** Under deadline pressure, it's easy to overlook a relevant exclusion or limitation that could have reduced or eliminated indemnity exposure. You find out when the file gets reviewed — or worse, in litigation.
- **No market context for coverage positions:** You know what the policy says, but you don't know how the provision compares to standard market language. That context strengthens your position — or tells you when a coverage argument is on shaky ground.
- **Outside counsel dependency:** Complex coverage questions get sent to outside counsel at \$500–\$800/hour when they could be resolved internally with the right analytical tools. Every referral is a budget line item and a cycle time delay.

Proof Points to Emphasize

- **Attorney-trained AI:** Built by a licensed coverage attorney who understands exactly how coverage determinations get challenged — and what it takes to build positions that hold up
- **Complex coverage analysis in hours, not days:** Instant access to all relevant policy provisions, exclusions, and endorsements for any claim scenario
- **Market intelligence:** Benchmark policy provisions against comparable programs to understand whether your coverage position is on solid or contested ground
- **User testimonial:** "After using it on layered programs and seeing it find discrepancies I missed, I was sold." — Claims professional at a top-15 brokerage

Note on Claims Personas:

Claims professionals sit on both sides of the table — carrier-side adjusters making coverage determinations and broker-side claims advocates arguing for coverage on behalf of their clients. Tailor the messaging accordingly:

- **Carrier-side adjusters:** Emphasize accuracy, defensibility, consistency, and loss ratio impact. They need to make the right call and document it.



- **Broker-side claims advocates:** Emphasize finding every basis for coverage, building the strongest position for the insured, and reducing reliance on outside counsel. They need to argue persuasively.

Both benefit from the same underlying capability — legal-grade coverage reasoning — but the framing should reflect their opposing objectives.

Integration Note for Claims:

Qumis works alongside your existing claims management system and email workflow. Coverage analysis can be initiated directly from your inbox and results integrated into your existing documentation process — no new system to learn, no disruption to your current workflow.

Suggested Talk Track Opening

When Competing Against Workflow Platforms:

"Workflow platforms excel at claims intake and data extraction — routing files and pulling information from forms. That's valuable for volume. But when you need to make an actual coverage determination — whether an exclusion applies, how endorsements interact, what your exposure is — that requires legal-grade policy reasoning, not document processing. Qumis was built by a coverage attorney specifically for these decisions. Many claims teams use workflow tools for intake and Qumis for the actual coverage analysis."

Cold Open for Claims Conversations:

"When a complex claim comes in, every hour matters — for the insured and for your exposure. Qumis instantly analyzes the policy and surfaces every relevant coverage provision, exclusion, and endorsement, then lets you benchmark those provisions against market standards. You get citation-backed coverage determinations in hours instead of days, with the documentation to defend your position. Let me show you how it works on a complex program."

5. Competitive Differentiation

The market for AI in commercial insurance is splitting into two categories: tools that automate insurance workflows and platforms that deliver coverage intelligence. Most vendors — regardless of funding or positioning — fall into the first category: they make existing processes faster without fundamentally changing the quality of coverage analysis. Qumis created and owns the second category. Our competitive frame must reinforce this distinction at every turn.



The Competitive Landscape

Qumis encounters four competitive alternatives in sales conversations, each requiring a distinct positioning approach:

1. Workflow automation platforms (FurtherAI, Fulcrum) — the most frequent head-to-head competitor
2. Generic AI tools (ChatGPT, Copilot, internal builds) — the most frequent budget objection
3. In-house coverage experts and outside counsel — the incumbent solution for complex coverage questions
4. Status quo (manual processes) — the "do nothing" alternative

vs. Workflow Automation Platforms (FurtherAI, Fulcrum)

They Say	We Say
"Automate your insurance workflows"	"Understand your coverage at a legal level — and benchmark it against the market"
Faster document processing	Deeper document understanding + market intelligence
Generic LLM with insurance UI	Purpose-built coverage reasoning engine trained by a licensed attorney
Speed and efficiency metrics	Accuracy, insight quality, and defensibility metrics
"Eliminate busywork"	"Elevate the quality of every coverage decision"
No proprietary data asset	Proprietary database of thousands of commercial programs

Competitive Talking Points:

FurtherAI (a16z-backed) and Fulcrum both raised large rounds in late 2025 competing on workflow automation. We compete on a fundamentally different dimension: legal-grade



coverage intelligence that understands what policies actually mean and benchmarks them against real market data. They make document processing faster. We make coverage decisions better.

Workflow platforms excel at submissions intake, data entry, and proposal generation — moving information between systems. Qumis excels at coverage reasoning — understanding policy language, identifying coverage gaps, explaining legal implications, and contextualizing coverage against market standards. Different problems, different solutions. Many of our customers use both: workflow tools for operational speed, Qumis for analytical depth. We are the intelligence layer that makes the entire technology stack smarter.

The complementary positioning is strategic, not defensive. Positioning Qumis alongside workflow tools rather than against them accomplishes three things: it neutralizes the "we already have FurtherAI" objection, it expands the budget conversation (different line item, different value), and it positions Qumis as the higher-value platform in the stack.

Killer Questions to Expose the Gap:

- "Can your AI explain why a pollution exclusion applies differently under occurrence vs. claims-made policies?" — Exposes the legal reasoning gap.
- "Can their platform tell you how that exclusion compares to what carriers are writing across comparable programs in your industry?" — Exposes the market intelligence gap. No workflow tool can answer this.
- "When your AI identifies a coverage gap, does it cite the specific policy language and explain the legal basis for the conclusion — or just flag a difference?" — Exposes the citation and defensibility gap.
- "Who on their founding team is a licensed coverage attorney?" — Exposes the domain expertise gap. This question is most effective with senior coverage professionals who understand what legal training means for policy interpretation.

Competitor-Specific Notes (internal use only — do not share with prospects):

- FurtherAI: Positions as "AI Workspace for Insurance" with workflow focus on submissions, audits, and claims intake. Founded by ML engineers. Strong on engineering talent and automation speed. Weak on coverage reasoning depth, legal defensibility, and market intelligence. They are building for operations teams; we are building for coverage professionals.
- Fulcrum: Positions as "Scale your accounts, not your headcount" with full lifecycle automation including AMS integration, proposals, and policy checking. Strong on



breadth of workflow coverage. Weak on analytical depth — their "policy checking" is surface-level compliance, not legal-grade coverage analysis with market context.

vs. Generic AI (ChatGPT, etc.)

This is the most common objection your sales team will encounter. It takes three forms, each requiring a different response:

Form 1: "Why can't we just use ChatGPT?" (The casual dismissal)

Response framework: Generic AI hallucinates on insurance questions because it has no domain-specific training, no insurance taxonomies, and no safeguards for when evidence is insufficient. It treats exclusions and conditions identically. It can't cite specific policy language. And it has zero proprietary market data to contextualize its output. Qumis was built by a licensed coverage attorney who encoded legal reasoning frameworks into the platform, backed by a proprietary database of thousands of commercial programs. The gap isn't small — it's structural.

Form 2: "Our IT team can build something with the API." (The build-vs-buy objection)

Response framework: Building a policy analysis tool on top of GPT is straightforward. Building one that delivers legal-grade accuracy with citation-linked outputs, multi-agent validation, insurance-specific taxonomies, and proprietary market benchmarking is a multi-year, multi-million-dollar effort that requires deep insurance domain expertise your IT team doesn't have. Every month you spend building, your competitors are already using Qumis. And you'll never replicate the proprietary data asset — that requires thousands of real commercial programs analyzed over years.

Form 3: "AI is improving fast — won't this be commoditized?" (The wait-and-see objection)

Response framework: Foundation models are improving rapidly — which makes the generic layer more commoditized, not less. What becomes more valuable as generic AI improves is the domain-specific intelligence layer: the legal reasoning frameworks, the insurance taxonomies, the proprietary market data, and the accumulated understanding of how policy language gets interpreted in practice. That's what Qumis is. We ride the improvement curve of foundation models while maintaining the domain expertise and data moat that generic tools can never replicate. The gap doesn't close over time — it widens.



Key differentiators (all forms):

- Citation-linked outputs: Every conclusion traces to source policy language. Generic AI provides answers without evidence.
- "I don't know" safeguards: When evidence is insufficient, Qumis says so. Generic AI confidently hallucinates — which in insurance creates E&O exposure.
- Insurance taxonomies: Qumis understands the structural roles of policy provisions (exclusions, conditions, endorsements, definitions). Generic AI treats them as undifferentiated text.
- Proprietary market data: Qumis benchmarks coverage against thousands of real commercial programs. Generic AI has no proprietary insurance data whatsoever.
- Audit trail: Qumis creates traceable, defensible documentation suitable for regulatory and litigation scrutiny. Generic AI outputs are ephemeral and untraceable.

Soundbite:

"You wouldn't use ChatGPT for a legal opinion. Why would you trust it with legal contracts worth millions in coverage?"

vs. Outside Counsel & Coverage Consultants

This competitor is often invisible in the sales process because it's not a software vendor — it's the law firm or coverage consultant the organization already uses for complex coverage questions. But it's the incumbent solution for exactly the use case where Qumis is strongest.

The Frame: Qumis doesn't replace outside counsel on genuinely complex disputes or litigation matters. But for the vast majority of coverage questions — the ones that currently generate a \$5,000–\$10,000 outside counsel invoice and a 2–3 week turnaround — Qumis delivers attorney-grade analysis in hours at a fraction of the cost. Reserve outside counsel for the matters that truly require it. Use Qumis for everything else.

Key Talking Points:

- **Economics:** Outside counsel bills \$500–\$800/hour for coverage analysis. A typical coverage opinion costs \$5,000–\$10,000 and takes 2–3 weeks. Qumis delivers comparable analysis in hours at a fraction of the cost.
- **Speed:** When a claim comes in or a renewal deadline approaches, waiting weeks for outside counsel isn't viable. Qumis provides immediate coverage reasoning that enables faster decisions.
- **Built by one of them:** Qumis was built by a licensed insurance coverage attorney who practiced coverage law. The analytical frameworks in the platform are the same ones coverage counsel use — encoded in software that operates at enterprise speed.
- **Scale:** Outside counsel doesn't scale. You can't send every coverage question to a law firm. Qumis makes legal-grade analysis available for every policy, every claim, every



renewal — not just the ones with enough dollars at stake to justify an outside counsel engagement.

- **Complementary, not replacement:** For matters heading to litigation or requiring formal legal opinions, outside counsel remains essential. Qumis handles the 80% of coverage questions that don't require a formal engagement — and provides better-prepared analysis for the 20% that do.

Soundbite: "Your outside counsel charges \$500/hour and takes three weeks. Qumis delivers the same analytical rigor in hours. Save outside counsel for the matters that truly need it."

vs. Status Quo (Manual Processes)

For prospects not yet evaluating AI solutions, the competitor is their current manual process — experienced professionals reading policies, comparing documents side by side, and relying on institutional knowledge.

The Frame: Manual coverage analysis works until it doesn't. It's slow, inconsistent across team members, and doesn't scale. Your best coverage people are bottlenecks, and when they leave, their expertise walks out the door. Qumis doesn't replace human judgment — it amplifies it, ensuring every team member operates with the same analytical depth and market context as your most experienced people.

Key Talking Points:

- **Consistency:** Manual processes produce different answers from different people on the same policy. Qumis standardizes coverage reasoning without eliminating human judgment.
- **Institutional knowledge risk:** Your best coverage people's expertise lives in their heads. Qumis captures and scales that analytical capability across the organization.
- **Scale:** Your team can't manually analyze every policy with the same rigor. Some get thorough reviews; others get skimmed. Qumis makes thorough analysis the default, not the exception.
- **Market context:** Even your most experienced people can only draw on the programs they've personally seen. Qumis benchmarks against thousands of commercial programs — broader market visibility than any individual can accumulate in a career.

Soundbite: "Your best coverage person has seen maybe a thousand programs in their career. Qumis has analyzed thousands — and brings that market context to every analysis, for every member of your team."

Competitive Positioning Summary



Competitor	Their Strength	Our Counter	Our Unassailable Advantage
Workflow Platforms	Operational speed, breadth of automation	We're complementary, not competitive — different value, different budget	Legal reasoning + proprietary market data. They can't build this.
Generic AI	Low cost, general capability, improving rapidly	Domain expertise, citation architecture, "I don't know" safeguards	Proprietary insurance data and attorney-trained reasoning. The gap widens as generic AI commoditizes.
Outside Counsel	Deep expertise, legal defensibility, litigation authority	Same analytical rigor at 10% of the cost and 10x the speed	Scale and speed. Law firms can't serve every coverage question economically.
Status Quo	Trusted, proven, no implementation risk	Amplifies human judgment, doesn't replace it; scales expertise across the team	Market intelligence from thousands of programs — broader than any individual's experience.

Rules of Engagement:

1. **Never disparage competitors by name to prospects.** Reference them by category ("workflow automation tools") unless the prospect brings up a specific vendor.
2. **Lead with the complementary frame** for workflow platforms. "We work alongside them" is stronger than "we're better than them" — it neutralizes the objection and elevates your positioning.
3. **Use killer questions sparingly and strategically.** They're most effective when a prospect is actively comparing you to a specific competitor, not as an opening move.
4. **Always bring it back to market intelligence.** This is the differentiator that ends competitive conversations. No competitor — workflow platform, generic AI, or law firm —



has a proprietary database of thousands of commercial insurance programs. When in doubt, land here.

5. **Competitor funding is internal context only.** Do not reference competitor funding amounts to prospects — it invites unfavorable comparisons. Your sales team should know the landscape; your prospects should hear about your value.

6. Proof Points & Social Proof

This section is the evidence layer of the entire messaging document. Every claim made in Sections 1–5 must be backed by something in this section. Organize proof points by category, tag each with the persona and section it supports, and note attribution requirements (named vs. anonymized). Update quarterly.

Enterprise Traction

1. **5 of the top 15 U.S. insurance brokerages** actively using Qumis in production
2. **127% net revenue retention** — customers are expanding, not just renewing. This is top-decile performance for seed-stage B2B SaaS.
3. **92% pilot-to-paid conversion rate** — every brokerage that has piloted Qumis has converted to a paid contract
4. **720% year-over-year ARR growth** (as of Q3 2025) — demonstrates accelerating product-market fit
5. **\$181K ARR in inbound pipeline generated in 6 weeks** (Q3 2025) — market pull, not just sales push

Use for: Brokerage executive messaging (Section 3), investor conversations, competitive positioning

Named Customer Logos

Customer	Type	Deployment Context
Brown & Brown	Top 10 US Brokerage	80+ active users across multiple offices
NFP	Top 15 US Brokerage	Expanded from 50-seat pilot to 250+ production users; deployed to middle-market team Q2 2025, expanded to construction practice Q4 2025
CRC Group	Wholesale Brokerage	Active deployment



Customer	Type	Deployment Context
Cottingham & Butler	Regional Brokerage	Active deployment
Kapnick	Regional Brokerage	Team of 20 using Qumis; 90% active retention rate; 718 average engagement score per user — highest per-user engagement of any account
IMA Financial	Top 25 US Brokerage	300+ policies analyzed in first 90 days
M3 Insurance	Regional Brokerage	Active deployment
Skyward Specialty	Specialty Carrier	Claims team deployment — source of key claims testimonials
Morris & Garritano	Regional Brokerage	100% active retention rate across 14 users — the "perfect customer profile"

Note: These are not logos on a wall. Each represents real production usage. When presenting logos, provide deployment context — it's more credible than names alone.

Use for: All sections. Diversify proof points by customer — avoid over-relying on any single logo.

Quantified Impact Metrics

Claims-Specific Metrics (sourced from customer interviews)

Metric	Value	Source	Attribution
Policy review speed	8x faster (hours → 5–10 minutes)	Brandon M., Senior Claims Advocate	Anonymized: "Senior Claims Advocate, Top 5 Regional Broker"
Claims response time	7x faster (~1 week → same/next day)	Brandon M.	Same attribution



Metric	Value	Source	Attribution
Claims capacity increase	+33% (150 → 200 claims/month)	Brandon M.	Same attribution
Additional claim recovery	\$1.8M in a single month	Brandon M.	Same attribution
Single claim save - denied claim (exclusion identification)	\$21,000	Vance Woods, Skyward Specialty	Named: "Senior Claims Adjuster, Skyward Specialty"
Claims resolution time	6 months → 45–60 days	Vance Woods	Same attribution
Initial claim review time	"couple hours" → 30 minutes	Vance Woods	Same attribution

Brokerage-Specific Metrics

Metric	Value	Source	Attribution
Policy review time reduction	80% faster (complex policy reviews)	NFP deployment	Named: "NFP"
Policy review time reduction	Cut in half; eliminated need for additional hire	Brian Pilarski, Kapnick	Named: "SVP & Partner, Kapnick"
Time savings (general)	5–7x on policy reviews	Multiple accounts	Anonymized: "consistently reported across accounts"
Claim save	\$1M on a challenged claim	Brown & Brown	Named: "Brown & Brown" — verify attribution permissions

Use for: Persona messaging (Section 4), competitive differentiation (Section 5, vs. status quo and outside counsel). Match metrics to the persona you're addressing — claims metrics for claims buyers, brokerage metrics for brokerage buyers.



Customer Testimonials

Coverage Expertise & Accuracy

Tom Hanekamp — Senior Partner, Cruser Mitchell (Coverage Counsel)

"I tested it with real-world coverage disputes, and it consistently delivered thoughtful, well-reasoned responses."

Context: Uses Qumis daily for coverage opinions. Also originated the "first-year associate" comparison and tested Qumis against ChatGPT on a racial discrimination exclusion analysis — Qumis got the coverage position correct where ChatGPT failed. Best for: Claims messaging, competitive differentiation vs. generic AI, credibility with coverage professionals

Connor Love — National Director of Growth, IMA Financial

"At first, I thought, 'There's no way this will catch everything.' But after using it on layered programs and seeing it track discrepancies I might have missed, I was sold."

Context: 300+ policies analyzed in first 90 days at IMA. Best for: Producer/front office messaging, skeptic conversion, layered program complexity proof point

Operational Efficiency & ROI

Brian Pilarski — SVP & Partner, Kapnick

"Qumis makes our policy reviews faster, sharper, and more accurate — it's a no-brainer. We cut our policy review time in half and eliminated the need for an additional hire."

Context: Team of 20 using Qumis since Q1 2025. Kapnick has the highest per-user engagement of any Qumis account (718 average engagement score). Best for: Account manager messaging, brokerage executive ROI conversations, headcount/efficiency proof point

Brandon M. — Senior Claims Advocate, Top 5 Broker (anonymized)

"Without Qumis, my caseload would be overwhelming. This isn't about working faster — it changed what's possible in my job."

Context: Source of the 8x faster, +33% capacity, and \$1.8M recovery metrics. Best for: Claims messaging, capacity/scale proof point, "not just faster — fundamentally different" positioning



Trust & Reliability

Vance Woods — Senior Claims Adjuster, Skyward Specialty

"What I would normally have to spend a couple hours on, I can condense into 30 minutes or less. It helps narrow down and logically note the file in terms of how you arrive at a decision or conclusion. The consistency's there — it's stable."

Context: Inland marine/cargo claims specialist. Saved \$21,000 on a single claim when Qumis identified a reefer exclusion buried in a 1,000-page policy that he had missed. The Vance story is one of the most compelling claims narratives — he was already negotiating a settlement at 70 cents on the dollar before Qumis caught the exclusion. Best for: Claims messaging, accuracy/defensibility proof point, the "what you're missing" argument

Joanne Snow — Senior Claims Technical Director, Skyward Specialty (D&O/Management Liability)

"I use it as a checklist to make sure I'm looking at everything I need to. It gives you a second pair of eyes — it confirms what you're thinking, or flags areas you need to look at. For these D&O claims with so many nuances, I think it's helpful for that."

Context: Experienced claims professional who positions Qumis as a validation tool, not a replacement. This framing is critical for credibility with senior coverage professionals who are skeptical of AI doing their job. Best for: Claims messaging to senior/skeptical practitioners, "second pair of eyes" positioning, D&O/specialty lines proof point

Daniel Berkowitz — Partner, Cruser Mitchell (Coverage Counsel)

Value articulation: "\$150/month vs. associate salary"

Context: Confirmed the first-year associate comparison with Hanekamp. The \$150/month vs. first-year associate salary framing is a powerful ROI anchor for coverage counsel conversations. Best for: Outside counsel ROI comparison, law firm sales conversations

Testimonial Selection Guide

Audience	Lead With	Support With
Brokerage executives	Brian Pilarski (Kapnick) — ROI + headcount	Connor Love (IMA) — skeptic conversion
Producers	Connor Love (IMA) — competitive edge	Brandon M. — capacity transformation



Audience	Lead With	Support With
Account managers	Brian Pilarski (Kapnick) — efficiency	Joanne Snow — accuracy validation
Claims leadership	Brandon M. — capacity + recovery metrics	Vance Woods — \$21K save + consistency
Claims practitioners	Vance Woods — day-to-day impact	Joanne Snow — "second pair of eyes"
Coverage counsel	Tom Hanekamp — legal credibility	Daniel Berkowitz — cost comparison
IT/Operations	Brian Pilarski (Kapnick) — team-wide adoption	NFP deployment scale

Credibility Markers

Founding Team

1. **Dan Schuleman, Esq.** — Co-Founder & CEO. Licensed insurance coverage attorney with law firm and in-house experience. This is the single most important credibility marker in the entire company — it backs the "attorney-trained AI" claim with a real person, not a marketing phrase.
2. **Shiv Sinha** — Co-Founder & CTO. Former Head of Application Development at Goldman Sachs; led development of Marcus, Goldman's consumer banking platform. Establishes technical credibility at enterprise scale.

Use for: All messaging. The attorney + enterprise tech founder combination is rare and should be referenced frequently. "Founded by a coverage attorney and the former head of application development at Goldman Sachs" is a one-sentence credibility bomb.

Institutional Validation

1. **SOC 2 Type I certified** — enterprise-grade security and compliance
2. **\$6.5M total funding** — \$2.2M pre-seed; \$4.3M oversubscribed seed round led by MTech Capital with American Family Ventures as new strategic investor and participation from all prior investors
3. **Selected for FinTech Innovation Lab New York** — top 3% of applicants. The program is run by the Partnership Fund for New York City and Accenture, with mentorship from senior executives at leading financial institutions.



4. **Oversubscribed pre-seed and seed rounds** — demand exceeded allocation, signaling strong investor conviction

Note on funding narrative: MTech Capital is an insurtech-focused fund; American Family Ventures is the venture arm of American Family Insurance, one of the largest P&C insurers in the US. The combination of specialist insurtech investors and strategic P&C capital validates both the technology thesis and the market opportunity.

Product & Technology

1. **Proprietary database of 10,000+ commercial insurance policies** across major lines of business — this is the data moat that competitors cannot replicate
2. **Multi-agent AI architecture** — specialized AI agents read, validate, and cross-reference documents, mirroring actual legal review processes
3. **Citation-linked outputs** — every conclusion traces to source policy language, creating auditable, defensible analysis
4. **"I don't know" safeguards** — the platform explicitly flags insufficient evidence rather than hallucinating, a critical trust differentiator in a profession where overconfident wrong answers create liability
5. **QumisScore** — proprietary coverage quality scoring that benchmarks policies against market standards

Case Studies (Narrative Format)

Case Study 1: "The \$21,000 Exclusion" — Vance Woods, Skyward Specialty

Situation: A cargo claim comes in. The insured's broker is calling every other day pushing for resolution. Vance's initial assessment: the claim is covered. He begins negotiating a settlement at 70 cents on the dollar.

What he missed: Buried in a 1,000-page policy was a reefer exclusion. The claim was based on temperature malfunction of the refrigeration unit — squarely within the exclusion. No coverage. But manually clicking through 1,000 pages under deadline pressure? He didn't catch it.

What Qumis found: When Vance ran the policy through Qumis, it identified the exclusion immediately. He generated a position letter citing the exact policy language. The broker's own investigation confirmed no coverage. Claim denied. File closed. \$21,000 saved on a single claim.

In his words: "The worst thing is a call from underwriting saying, 'Did you not see that exclusion?' After a while, the company will tell you to take your equipment and go."



Best for: Claims sales conversations, live demo narrative, conference presentations. This story is visceral — it's about job security, not just efficiency.

Case Study 2: "The Capacity Transformation" — Brandon M., Top 5 Regional Broker

Situation: Senior claims advocate managing hundreds of active claims across complex verticals at a top-5 US brokerage. Caseload is at capacity. Adding staff would take months and \$150K+ in salary.

Impact with Qumis:

1. Policy review time: 8x faster (hours → 5–10 minutes)
2. Claims response time: 7x faster (~1 week → same/next day)
3. Claims capacity: +33% (150 → 200 claims/month)
4. Additional claim recovery: \$1.8M in a single month
5. Headcount impact: Eliminated need for additional hire

In his words: "Without Qumis, my caseload would be overwhelming. This isn't about working faster — it changed what's possible in my job."

Best for: Claims leadership ROI conversations, board presentations, investor deck. The \$1.8M recovery number is the single most powerful ROI data point in the entire Qumis story.

Case Study 3: "Land and Expand" — NFP

Situation: NFP, a top-15 US brokerage, began with a 50-seat pilot in Q2 2025 focused on the middle-market team.

Expansion trajectory: Within months, the deployment expanded to 250+ production users and extended into NFP's construction practice group in Q4 2025. This represents the land-and-expand motion that drives 127% NRR.

What this proves: Enterprise insurance organizations don't just try Qumis — they deploy it broadly once they see results. The expansion from middle-market to construction practice demonstrates cross-practice-group utility and internal champion advocacy.

Best for: Brokerage executive conversations ("this is what adoption looks like at organizations your size"), investor deck, competitive positioning.



7. Language Guide

This section governs how Qumis talks about itself — in marketing copy, sales conversations, investor materials, customer communications, and internal documents. Language discipline is what turns positioning into brand. If a term isn't here, think twice before using it. If a term is in the "Avoid" list, don't use it regardless of context.

Core Vocabulary: Terms to Use

Term	Why It Works	Primary Context
Coverage intelligence	Category-defining. Positions Qumis above "automation" and "document processing." This is the term that should appear in every headline, every first paragraph, every elevator pitch.	All contexts — this is the category name
Legal-grade	Signals precision and defensibility without claiming to be legal advice. The hyphen matters — it's a compound modifier, not a noun. Always "legal-grade analysis," never "legal analysis."	Marketing, sales, investor materials
Attorney-trained AI	Credibility anchor. Differentiates from generic AI and workflow tools. Implies a specific person (Dan) built the reasoning frameworks, not just that lawyers were "consulted."	Marketing headlines, competitive positioning, first-meeting sales conversations
Citation-linked	Demonstrates accountability. Directly addresses the hallucination concern that every informed buyer has. More specific and credible than "accurate" or "reliable."	Product descriptions, competitive differentiation vs. generic AI, IT/compliance conversations
Coverage reasoning	Conveys analytical depth — Qumis doesn't just find text, it interprets meaning. Distinguishes from search/retrieval tools.	Product descriptions, technical conversations, persona messaging for coverage professionals



Market intelligence	Signals the proprietary data asset — Qumis doesn't just analyze a single policy, it contextualizes coverage against thousands of real commercial programs. This is the term that ends competitive conversations.	Sales conversations (especially later-stage), competitive differentiation, investor materials
Proprietary data / Proprietary database	Establishes the data moat. "Proprietary" signals that this asset cannot be replicated by competitors. Always pair with scale: "proprietary database of thousands of commercial insurance programs."	Competitive differentiation, investor materials, enterprise sales
Coverage benchmarking	Translates the data moat into a buyer benefit. Practitioners understand benchmarking — it's how they already think about market positioning.	Producer messaging, account manager messaging, executive conversations
Enterprise-grade	Signals security, scale, and seriousness. Use in IT/operations contexts.	IT buyer conversations, security documentation, RFP responses
Built by experts, for experts	Establishes peer credibility with practitioners who are skeptical of "tech people" building insurance tools.	Marketing, first-meeting positioning, conference presentations
Coverage intelligence platform	The full product descriptor. Use when you need to name what Qumis is (not just what it does). Replaces "AI platform for insurance" or "policy analysis tool."	Website, press releases, formal descriptions
Insurance-native	Signals that Qumis was built specifically for insurance from the ground up — not a generic AI tool with an insurance skin.	Competitive differentiation, product descriptions
QumisScore	Proprietary coverage quality scoring. Use when discussing the market intelligence capability specifically. Treat as a proper noun (capitalized, one word).	Product demos, account manager messaging, expansion conversations

Terms to Avoid



Avoid	Why	Say Instead
"AI-powered" (alone)	Generic. Every company says this. It communicates nothing about what makes Qumis different.	"Attorney-trained AI" or "insurance-native AI"
"Game-changer"	Clichéd marketing-speak. Let customers say it — and they do (see testimonials). When Qumis says it, it undermines credibility with skeptical insurance buyers.	Let the proof points speak. Or: "fundamentally different approach"
"Automate" / "Automation" (as a lead or primary descriptor)	Cedes positioning to workflow tools. If Qumis is "automation," buyers compare you to FurtherAI and Fulcrum on their terms.	"Intelligence" or "analysis" as lead terms. Exception: it's acceptable to say "Qumis adds intelligence to your existing automation" in the complementary frame — but Qumis itself is never the automation.
"GPT" / "ChatGPT" (in Qumis descriptions)	Anchors to generic AI and invites direct comparison. Only reference ChatGPT in competitive differentiation contexts ("You wouldn't use ChatGPT for legal advice"), never in Qumis product descriptions.	"Purpose-built AI" or "coverage reasoning engine"
"Revolutionary" / "Disruptive"	Oversold. Insurance buyers have heard this from every vendor for 20 years. These words signal "startup that doesn't understand our industry."	Show, don't tell. Let the metrics and testimonials make the case.
"Effortless" / "Seamless"	Undermines the expertise of your users. Coverage analysis is serious, consequential work. Calling it "effortless" disrespects the profession.	"Faster" or "more efficient" — acknowledge the work while showing how Qumis elevates it



"Copilot" (unqualified)	Overused across every AI product. Diluted to meaninglessness.	If used at all, always qualify: "coverage copilot." But prefer "coverage intelligence platform."
"Solution"	Empty corporate-speak. A "solution" could be anything.	Be specific: "platform," "analysis," "intelligence," or name the actual capability
"Leverage" / "Utilize"	Corporate jargon that makes copy feel generic.	"Use," "apply," "deploy" — plain language is more credible with practitioners
"End-to-end"	Implies workflow automation breadth, which is the competitor's positioning. Qumis wins on depth, not breadth.	Describe the specific capability: "from initial analysis through defensible documentation"
"Policy analysis tool"	Undersells the platform. A "tool" is commoditized; a "platform" is infrastructure. "Analysis" alone misses the market intelligence dimension.	"Coverage intelligence platform"

Key Phrases by Context

Strategic Positioning Phrases (use in marketing, investor materials, executive conversations)

"Understanding the policy is the hardest problem in insurance." — *This is the single most important strategic phrase in Qumis's vocabulary.* It reframes the entire market conversation from workflow speed to analytical depth. Use it to open conversations, open presentations, and anchor competitive discussions. When a buyer agrees with this statement, they've already self-selected into Qumis's value proposition.

"Coverage intelligence, not workflow automation." — The category distinction in six words. Use when positioning against competitors or when a buyer tries to lump Qumis with workflow tools.

"Law-firm accuracy at enterprise scale." — Concise, memorable, captures both the legal expertise and the scalability. Strong headline language.



"We deliver not just what a policy says, but what it means in the context of the broader market." — The dual value proposition in one sentence. Use when explaining the positioning to someone for the first time.

Sales Conversation Phrases (use in demos, discovery calls, follow-up emails)

"Like having a coverage attorney on every account." — The evolved version of the first-year associate comparison. Works for brokerage executives and anyone who understands the cost of outside counsel. Implies scale (every account) and quality (attorney-level).

"Like having a first-year associate — only faster and more accurate." — Tom Hanekamp's original framing. Most effective with claims professionals, coverage counsel, and anyone who has managed junior staff. Use when the audience has legal context.

"Your most experienced person has seen maybe a thousand programs in their career. Qumis has analyzed thousands — and brings that market context to every analysis." — The market intelligence proof point. Use when differentiating against the status quo or when a buyer says "our people are good enough."

"Qumis is the intelligence layer that makes your entire technology stack smarter." — The complementary positioning frame. Use when the buyer already has FurtherAI, Fulcrum, or another workflow tool and you need to position alongside rather than against.

"We built the 'I don't know' into the system on purpose." — Use when addressing hallucination concerns. This flips a perceived AI weakness into a trust differentiator. Particularly effective with compliance-oriented buyers and IT decision-makers.

Competitive Phrases (use when directly compared to alternatives)

"You wouldn't use ChatGPT for a legal opinion. Why would you trust it with legal contracts worth millions in coverage?" — The generic AI soundbite. Strong opener, but be prepared for the follow-up objection: "We're not using it for legal opinions, just for faster document review." Response: "That's exactly the risk — when a tool doesn't know the difference between document review and coverage analysis, it can't tell you when the answer matters."

"Can their platform tell you how that exclusion compares to what carriers are writing across comparable programs?" — The killer question that exposes the market intelligence gap. No workflow tool or generic AI can answer this.

"Who on their founding team is a licensed coverage attorney?" — The credibility question. Most effective with senior coverage professionals. Use sparingly and strategically — it's a knockout punch, not a jab.



"Different problems, different solutions. Many of our customers use both." — The complementary frame for workflow competitors. Neutralizes the "we already have X" objection.

Customer Success Phrases (use in onboarding, QBRs, renewal conversations)

"The goal isn't to replace your judgment — it's to make sure every member of your team operates with the same analytical depth as your best people." — The amplification frame for expansion conversations.

"How many coverage questions does your team handle that currently don't get a thorough analysis because there isn't enough time?" — The unmet demand question. Reveals the addressable opportunity within the account and supports expansion.

Competitor Reference Language

When referring to competitors in any context, use these category terms consistently:

Competitor	How to Reference (External)	How to Reference (Internal)
FurtherAI	"Workflow automation tools"	FurtherAI (by name, internal only)
Fulcrum	"Workflow automation platforms"	Fulcrum (by name, internal only)
ChatGPT / Copilot	"Generic AI" or "general-purpose AI"	ChatGPT, Copilot (by name)
Internal builds	"Custom-built solutions"	"Build vs. buy" frame
Outside counsel	"Outside counsel" or "coverage consultants"	Same
Manual processes	"Current manual process" or "status quo"	Same

Rule: Never reference a competitor by name to a prospect unless the prospect brings them up first. Always refer by category. If the prospect names a specific competitor, you may respond using the competitor's name — but redirect to category framing as quickly as possible.



Regulatory & Legal Guardrails

"Legal-grade" is Qumis's most powerful term. It is also the term most likely to create liability if misused. These guardrails are non-negotiable:

- **Qumis does not provide legal advice.** "Legal-grade" refers to the quality, rigor, and methodology of the analysis — not to the legal authority of the output. Every external-facing description of Qumis should be reviewed to ensure it does not imply that Qumis output constitutes a legal opinion.
- **"Attorney-trained" means the platform's reasoning frameworks were built by a licensed coverage attorney.** It does not mean every output has been reviewed by an attorney.
- **All marketing copy, website pages, and sales collateral should include appropriate disclaimers** where Qumis capabilities are described in terms of legal analysis. Standard disclaimer language: "Qumis provides coverage intelligence and analysis tools. Qumis output does not constitute legal advice. Users should consult qualified legal counsel for formal legal opinions."
- **In sales conversations**, if a prospect asks "Is this legal advice?" the answer is: "No — Qumis provides attorney-grade coverage analysis that supports your team's decision-making. For matters requiring a formal legal opinion, we recommend consulting qualified counsel. Many of our customers use Qumis to handle the 80% of coverage questions that don't require outside counsel, and to better prepare for the 20% that do."

Tone Principles

These govern *how* Qumis communicates, not just *what* it says:

1. **Confident, not arrogant.** Qumis has earned the right to be assertive about its capabilities — real customers, real metrics, real differentiation. But insurance is a relationship business built on trust. Overconfident claims alienate the exact buyers you're targeting. Let proof points carry the weight.
2. **Precise, not technical.** Insurance professionals respect precision in language — it's literally their profession. Vague claims ("better," "smarter," "faster") without specifics feel like marketing. Specific claims ("8x faster policy review," "citation-linked to source policy language") feel like someone who understands the work.
3. **Respectful of the profession.** Insurance coverage professionals are experts. They've spent decades building judgment that AI cannot fully replicate. Qumis messaging should never imply their work is simple, rote, or easily automated. The frame is always amplification, not replacement.
4. **Domain-fluent.** Use insurance terminology correctly. "Insured," not "customer." "Exclusion," not "carve-out" (unless in the correct technical context). "Endorsement," not "rider" (unless discussing personal lines). Getting the language right signals that Qumis was built by someone who practiced in this space — because it was.



5. **Direct over clever.** Insurance professionals are skeptical of marketing polish. Plain, direct language ("Qumis finds coverage gaps your team might miss") outperforms clever copywriting ("The policy is both the starting point and the sticking point") every time.

8. Objection Handling

This section arms your team for the hardest moments in sales conversations. Each objection includes the initial response, the most likely follow-up, and the response to that follow-up. Objections rarely end after one exchange — prepare for the conversation, not just the question.

Principles for all objection handling:

1. Acknowledge before countering. Dismissing a prospect's concern makes them defensive. Validating it builds trust.
2. Lead with Qumis, not competitors. Spend more time on what you are than what they aren't.
3. Land on market intelligence when possible. The proprietary data asset is the argument that ends objection conversations.
4. Use proof points. Abstract claims are forgettable. Specific customer stories are persuasive.

Objection 1: "How is this different from ChatGPT?"

This is the most common objection. It takes multiple forms depending on who's asking.

Initial Response: Qumis is a purpose-built coverage intelligence platform — built by a licensed coverage attorney, trained on insurance-specific reasoning frameworks, and backed by a proprietary database of thousands of real commercial insurance programs. Every output is citation-linked to the source policy language, so you can verify any conclusion. And when the evidence is ambiguous or insufficient, Qumis tells you — it says "I don't know" rather than guessing. That's not how ChatGPT works. Generic AI has no insurance domain expertise, no proprietary data, no citation architecture, and no safeguards against confident wrong answers. In insurance, a confident wrong answer creates liability.

Follow-up: "But ChatGPT is free / much cheaper. Isn't it good enough for most tasks?"

For basic document search, maybe. But coverage analysis isn't basic document search — it requires understanding how exclusions interact with definitions, how endorsements modify coverage, and how sublimits affect recovery. ChatGPT treats a policy as undifferentiated text. Qumis understands the structural roles of every provision. And here's the part generic AI can never replicate: Qumis benchmarks coverage against thousands of real commercial programs. When Qumis identifies a coverage gap, it can tell you how that gap compares to what carriers



are writing across comparable programs in your industry. ChatGPT has no proprietary insurance data whatsoever.

Follow-up: "AI is improving fast. Won't GPT-5 or GPT-6 solve this?"

Foundation models are improving rapidly — which actually makes our position stronger, not weaker. As generic AI gets better, the *generic* layer commoditizes. What becomes more valuable is the domain-specific intelligence layer: the legal reasoning frameworks, the insurance taxonomies, the proprietary market data accumulated from thousands of real programs, and the understanding of how policy language gets interpreted in practice. That's what Qumis is. We ride the improvement curve of foundation models while maintaining the domain expertise and data moat that generic tools can never replicate. The gap doesn't close over time — it widens.

Objection 2: "We already have [workflow automation tool]"

Initial Response: That's actually a great foundation. Workflow tools like that are strong at operational efficiency — submissions intake, data movement, document processing. Qumis does something different: we deliver coverage intelligence — understanding what the policy actually means, identifying coverage gaps, and benchmarking coverage against market standards. Think of it this way: your workflow tool can tell you two policies are different. Qumis tells you what those differences mean for coverage and how they compare to what the rest of the market is writing. Many of our customers use both — the workflow tool for operational speed, Qumis for analytical depth. We're the intelligence layer that makes the entire stack smarter.

Follow-up: "That sounds like we'd be paying for two tools that overlap."

Different budget, different value, different users. Your workflow tool serves operations teams processing high volumes of documents. Qumis serves coverage professionals — the people who need to understand, interpret, and defend policy positions. There's almost no user overlap. And the ROI model is completely different: workflow tools save time on administrative tasks. Qumis prevents coverage errors, identifies gaps your team might miss, and provides the market context that supports better coverage decisions. One of our customers — a senior claims advocate at a top-5 broker — recovered an additional \$1.8M in claims in a single month using Qumis. That's not an efficiency gain. That's a revenue outcome no workflow tool delivers.

Objection 3: "How do I know I can trust the AI output?"

Initial Response: This is exactly the right question, and it's why we built Qumis the way we did. Three things make Qumis output trustworthy in a way generic AI is not. First, every conclusion is citation-linked to the source policy language — you click any finding and see exactly which provision it came from. Second, we built "I don't know" safeguards into the system on purpose. When the evidence is ambiguous or a document is unclear, Qumis flags it rather than guessing.



In insurance, a confidently wrong answer is worse than no answer at all. Third, Qumis creates an auditable documentation trail — analysis that's traceable, defensible, and suitable for regulatory or litigation scrutiny.

Follow-up: "But how accurate is it really? What's your error rate?"

We don't publish a single accuracy number because accuracy varies by policy type, complexity, and use case — and anyone who gives you a single percentage is oversimplifying. What we can tell you is this: one of our users — a senior claims adjuster at a specialty carrier — told us Qumis caught a reefer exclusion buried in a 1,000-page policy that he had missed. He was already negotiating a settlement. That single catch saved \$21,000 on one claim. And a senior claims technical director at the same carrier described Qumis as a "second pair of eyes" that "confirms what you're thinking, or flags areas you need to look at." The trust architecture isn't about a percentage — it's about citations, safeguards, and the ability to verify everything the system tells you.

Objection 4: "Is this going to replace my team?"

Initial Response: No — and we designed it that way deliberately. Qumis is a force multiplier for your existing team, not a replacement. Your senior coverage professionals still make the decisions, still exercise judgment, still manage client relationships. What Qumis does is make sure every member of your team operates with the same analytical depth and market context as your most experienced people. Think of it as giving each person on your team a tireless coverage expert who has already read the entire policy and analyzed it against thousands of comparable programs — so your people can focus on the judgment calls that actually require their expertise.

Follow-up: "My team is worried about this. How do I get buy-in?"

The practitioners who are most skeptical upfront often become the strongest advocates. One of our users, Joanne Snow — a Senior Claims Technical Director at Skyward Specialty who handles complex D&O and management liability claims — put it this way: "I use it as a checklist to make sure I'm looking at everything I need to. It gives you a second pair of eyes — confirms what you're thinking, or flags areas you need to look at." That's how experienced professionals actually use Qumis — not as a replacement for their expertise, but as a tool that makes their expertise more thorough and consistent. The best way to build buy-in is a focused pilot with your most demanding team members. When your toughest critics see Qumis catch something they missed on their own policies, adoption follows naturally.

Objection 5: "What about security and compliance?"

Initial Response: We're SOC 2 Type I certified and built specifically for enterprise insurance deployment. Three things that matter most for your security review: client data is never used for



AI model training — your policies stay your policies. We maintain strict data isolation between tenants — no client can access another client's data. And we provide enterprise-grade access controls including SSO and MFA. Multiple top 15 U.S. brokerages put us through their full security review process and approved us for production use with their most sensitive policy documents.

Follow-up: "We need to run this through our InfoSec team. What does that look like?"

We've been through rigorous security reviews with major enterprise brokerages and can provide your InfoSec team with our SOC 2 report, security questionnaire responses, and architecture documentation. Typical security reviews take two to four weeks. We're happy to schedule a technical call with your security team directly — our CTO has led enterprise security implementations at Goldman Sachs and can speak their language.

Objection 6: "How long does implementation take?"

Initial Response: Most teams are productive within days. Qumis is cloud-based — no on-premise installation, no complex integration required. You upload policies and start analyzing them. There's no heavy IT lift, which is why our typical implementation timeline is days to weeks, not months. For teams wanting deeper integration with existing systems — AMS platforms like Applied Epic or Sagitta, document management systems, email workflows — we have APIs and can scope a phased rollout.

Follow-up: "What does the rollout typically look like?"

Most customers start with a focused deployment — a specific practice group, a specific use case like renewals or claims analysis — and expand as they see results. That's how NFP went from a 50-seat pilot focused on their middle-market team to 250+ production users across multiple practice groups including construction. We provide hands-on onboarding, training tailored to your team's workflows, and ongoing customer success support. The goal is time to value in the first week, measurable impact in the first month.

Objection 7: "We tried [other AI tool] and it didn't work"

Initial Response: That's a common experience, and frankly it's why many of our best customers came to us — they were burned by generic AI that doesn't understand insurance. Most AI tools applied to insurance are either general-purpose models with an insurance interface or workflow automation tools that don't attempt real coverage analysis. Qumis is neither. Our founder is a licensed coverage attorney who experienced the failures of generic AI on real coverage questions and built Qumis specifically to solve what those tools couldn't. The reasoning frameworks in the platform are the same ones coverage counsel use — encoded in software, backed by a proprietary database of thousands of real commercial programs.



Follow-up: "How do I know this time will be different?"

Two things. First, let us show you on your actual policies. We'll run a pilot on your documents, not a demo on ours. When you see Qumis accurately analyze a layered program or identify a coverage gap in a policy you know well, you'll see the difference immediately. Second, ask us the hard questions — not simple policy lookups, but the nuanced coverage questions that tripped up whatever tool you used before. We had a user — Connor Love, National Director of Growth at IMA — who told us he initially thought "There's no way this will catch everything." After using it on layered programs and seeing it track discrepancies he might have missed, he was sold. Let your skeptics put us to the test. That's how we win.

Objection 8: "This is too expensive" / "We don't have the budget"

Initial Response: I understand — let me reframe what you're actually comparing against. The alternative to Qumis isn't free. It's the cost of the manual process: your most experienced people spending hours on policy reviews that Qumis handles in minutes. It's the coverage gap that gets missed because there wasn't time for a thorough review. It's the outside counsel engagement at \$500 to \$800 an hour for a coverage question Qumis could have answered immediately. One of our customers, Brian Pilarski at Kapnick, told us Qumis eliminated the need for an additional hire while making their reviews faster and more accurate. That single hire would have cost significantly more than Qumis across a full team.

Follow-up: "Can we start smaller?"

Absolutely — that's exactly how most of our enterprise customers deploy. Start with a focused team of 10 to 20 users on a specific use case. Measure the impact over 90 days. Then expand based on results. That's the motion that drove NFP from 50 seats to 250+ and gives us 127% net revenue retention across our customer base. We're confident enough in the value that we want you to prove it to yourself before scaling.

Objection 9: "Our IT team can build this internally"

Initial Response: Building a document chat interface on top of a foundation model is straightforward — your team could have that running in weeks. But that's not what Qumis is. Building a platform that delivers legal-grade coverage analysis with citation-linked outputs, multi-agent validation, insurance-specific taxonomies, "I don't know" safeguards, and proprietary market benchmarking drawn from thousands of real commercial programs — that's a multi-year, multi-million-dollar effort that requires deep insurance domain expertise your engineering team doesn't have. And there's one thing you can never build internally: the proprietary data asset. That requires thousands of real commercial programs analyzed over years of enterprise customer relationships.

Follow-up: "We have insurance domain experts who could guide the build."



Domain experts are critical — and they should be the ones using Qumis, not spending their time as product managers for an internal AI project. Every month your team spends building, your competitors are already using coverage intelligence to win accounts and identify gaps. The question isn't whether your team *could* build something — it's whether that's the highest-value use of your most expensive people's time, and whether you can reach the level of accuracy, market data, and coverage reasoning that Qumis delivers today.

Objection 10: "We need to see it work on our policies first"

Note: This is a buying signal, not a true objection. Treat it as forward momentum.

Response: That's exactly how we prefer to start. Generic demos on sample policies don't prove anything. Let's set up a focused pilot on your actual documents — the policies your team knows best — so you can evaluate the analysis against your own expertise. Most pilots run two to four weeks with a defined scope: a specific practice group, a specific line of business, or a specific use case like renewals or claims analysis. Our customer success team provides hands-on onboarding and is available throughout. When your team sees Qumis accurately analyze the policies they've spent years learning, the conversation changes.

Quick Reference: Objection → Proof Point Pairing

Objection	Best Proof Point
"How is this different from ChatGPT?"	Tom Hanekamp's ChatGPT vs. Qumis test (racial discrimination exclusion — Qumis got the coverage position correct, ChatGPT failed)
"We already have a workflow tool"	\$1.8M additional claim recovery (Brandon M.) — an outcome no workflow tool delivers
"How do I trust the output?"	Vance Woods \$21K save — Qumis caught what he missed on a 1,000-page policy
"Will it replace my team?"	Joanne Snow — "second pair of eyes" framing from a senior practitioner



"Security concerns" 5 of top 15 US brokerages passed us through full security review

"Implementation timeline" NFP: productive in days, expanded from 50 to 250+ users

"We tried AI before" Connor Love (IMA) — initial skeptic, sold after seeing results on layered programs

"Too expensive" Brian Pilarski (Kapnick) — eliminated need for additional hire

"We can build internally" Proprietary database of thousands of programs — years of accumulated data you can't replicate

"Need to see it on our policies" 100% pilot-to-paid conversion rate — every company that's tried it has bought it

9. The Qumis Ascend Program

Ascend is Qumis's evidence-backed implementation and change management program. It is a core differentiator in enterprise sales conversations — not a professional services add-on, but a built-in component of every enterprise engagement. This section covers why Ascend exists, how to position it, and why no competitor can replicate it.

Why Ascend Exists: The Evidence Base

Every element of Ascend traces back to real deployment data, not borrowed consulting methodology. Three findings from our own platform data and enterprise deployment experience shaped the program:

The activation threshold. Users who analyze 20 or more documents within their first 30 days retain at 81%. Users below that threshold retain at 36% — a 45-percentage-point gap. Ascend is designed to drive every user past that threshold within the critical first month.



The ghost user problem. When we scaled enterprise accounts without structured change management, ghost rates — users who received licenses but never activated — increased sevenfold. At one account, nearly half of provisioned users never logged in. The technology was identical. What was missing was structured onboarding, internal champion development, and workflow-specific use case design. Ascend was built to prevent this.

The Center of Excellence pattern. Our highest-performing accounts (90–100% active rates, 3–5x higher per-user engagement) share a common structure: they have internal "architects" — experienced coverage professionals who define how the team uses Qumis and serve as the go-to resource for colleagues. Ascend systematically creates this pattern at every deployment rather than hoping it emerges organically.

How to Position Ascend

Ascend is not "onboarding." It is a structured, phased deployment program — discovery, pilot, validated expansion, ongoing optimization — designed by insurance professionals who understand coverage workflows at a practitioner level. The core message: **Qumis has deployed coverage intelligence at five of the top 15 U.S. brokerages and learned exactly what drives adoption and what kills it. Ascend is those lessons encoded into a repeatable program backed by real data.**

The framing is always "this is how we ensure you succeed" — never "this is an extra service we offer." Ascend should appear in the sales process during discovery, not at contracting. When a prospect asks "what does implementation look like?", Ascend is the answer.

To brokerage executives and operations leaders:

Ascend is how you de-risk an AI deployment. Most AI implementations in insurance fail not because the technology doesn't work, but because the change management isn't there. Our structured deployments achieve 90%+ active utilization rates, compared to the 47% industry average for enterprise SaaS license utilization. We built Ascend because we've seen both sides — we know what works and what doesn't, and we don't succeed unless your team actually uses the platform.

To IT and operations decision-makers:

Ascend minimizes the IT lift and maximizes adoption speed. We handle change management, role-specific training, and ongoing support so your team doesn't have to build an internal AI enablement function from scratch. Typical timeline: productive users within days, measurable impact within 30 days, validated expansion within 90 days.

To department heads and team leaders:



Ascend means your team isn't left to figure it out on their own. We build custom prompt libraries and analytical frameworks specific to your lines of business, provide hands-on training using your actual policies, and develop internal champions on your team who become the go-to experts — so adoption sustains long after formal onboarding ends.

To skeptical practitioners:

This is exactly why we built Ascend. We've seen what happens when AI tools get dropped on insurance professionals without proper context — they try it once, get confused, and never come back. Ascend is different because it's designed by people who've done coverage work professionally. The training uses your actual policies, not generic demos. The prompt templates reflect how coverage professionals actually think. And we track activation in real time — if someone is stuck, we know within days and intervene before they disengage.

Ascend as Competitive Differentiator

Ascend is a competitive moat that workflow automation vendors and generic AI tools cannot replicate because it requires three things they don't have:

Insurance domain expertise. The people designing and delivering Ascend understand coverage analysis at a professional level. They can build prompt libraries that reflect how a D&O specialist thinks differently from a construction practice leader. They can facilitate training using real coverage questions and diagnose whether low adoption is a product issue, a workflow fit issue, or a training gap. Generic AI vendors can train users on their interface. They cannot train users on coverage analysis.

Proprietary deployment data. The activation thresholds, engagement benchmarks, champion patterns, and expansion triggers come from real usage data across hundreds of enterprise insurance users. This is evidence-based activation science specific to AI adoption in insurance, not theoretical change management methodology.

A proven expansion playbook. The phased approach — pilot to validated expansion to full deployment — has been refined through engagements that grew from initial pilots to hundreds of production users. The expansion path from 50 seats to 250+ that we've executed at enterprise accounts happened because the Ascend framework systematically creates the conditions for growth: executive sponsorship, internal champions, quantified impact, and a clear business case for the next phase.

When a prospect asks "what makes your implementation different?", the answer is: our competitors can provision accounts. They can't do this.



Ascend Proof Points

Metric	Data Point
Activation threshold	81% retention for users reaching 20+ documents in 30 days vs. 36% below
Active utilization (structured)	90%+ at Kapnick, Morris & Garritano vs. 47% industry SaaS average
Ghost rate improvement	<5% at structured deployments vs. 29% when scaling without Ascend
Pilot conversion	100% of companies completing a structured pilot become paying customers
Expansion velocity	50 → 250+ users at NFP following phased model
Per-user engagement	383–718 at structured deployments vs. 126 at unstructured
Time to productivity	Days, not months

Key Phrases for Ascend Conversations

"We've deployed at five of the top 15 U.S. brokerages. Ascend is everything we learned about what drives adoption — encoded into a program."

"Our structured deployments achieve 90%+ utilization. The industry average for enterprise SaaS is 47%. The difference is change management."



"We don't just give your team a login. We build the workflows, develop the champions, and stay engaged until adoption is self-sustaining."

"Whatever we do, it's got to be a contained, systematic rollout that shows we know what we're doing and we've done it before." (*Internal framing — captures the Ascend philosophy.*)

10. Data, AI, and Customer Privacy: Messaging Guide

This section governs how Qumis talks about customer data and AI. Every claim must be accurate to the actual architecture and contractual commitments. The language is precise for a reason — imprecise data messaging either creates legal liability or forfeits competitive positioning.

The Core Tension

Qumis's positioning depends on two commitments that sound contradictory:

1. "Your data is never used to train large language models."
2. "Qumis has proprietary intelligence built from thousands of commercial insurance programs."

These are not contradictory. They describe different activities happening to different forms of data. The first refers to customer-identifiable data, which is always private and siloed. The second refers to de-identified, aggregated coverage intelligence — patterns, structures, benchmarks — that cannot be traced back to any customer. Your team must be able to explain this distinction clearly. When they can't, trust collapses.

The precise commitment: Qumis does not use customer data to train general-purpose or foundational large language models. Qumis does build proprietary knowledge infrastructure and specialized domain capabilities from de-identified, aggregated coverage intelligence. These are architecturally and contractually distinct activities. The prohibition is specific. The permitted uses are specific. Both are in the MSA (Section 8.5 and Section 4.4).

What Your Team Needs to Understand

Customer-identifiable data — policies, analysis outputs, usage history — is logically segregated, never visible to other customers, never used to train any model, and deleted within 60 days of termination.

Third-party AI providers process customer data for inference only. They are contractually prohibited from training on it.



De-identified, aggregated intelligence — coverage language patterns, carrier behaviors, endorsement structures, market benchmarks — is derived from customer data but stripped of identifying information. It powers Qumis's market intelligence features and is used to improve the platform, including developing specialized domain capabilities like classifiers, quality scoring, and analytical tools. This is standard practice for vertical AI platforms and is analogous to how a credit bureau derives market benchmarks from individual account data.

User feedback signals (thumbs up/down, corrections) inform product improvement. They are about output quality, not customer documents.

The Questions Your Team Will Get

"Do you use our data to train your AI?"

"We do not use your data to train large language models — not ours, not any third-party provider's. That's a contractual commitment. Separately, Qumis maintains a proprietary knowledge base built from de-identified coverage patterns across thousands of programs — that's what powers our market intelligence and benchmarking. But that's aggregated and de-identified. It doesn't contain anyone's specific policies."

"How can you have market intelligence if you don't use customer data?"

"We distinguish between customer data — always private, always siloed — and de-identified coverage intelligence derived from thousands of programs. The knowledge base captures patterns: how carriers structure exclusions, what coverage terms look like across industries. It never contains specific policies or identifying information. It's like how a credit bureau provides market benchmarks without exposing any individual's account."

"Is our data commingled with other customers'?"

"No. Your data is logically segregated — no other customer can access or view it. Our market intelligence layer is a separate, de-identified knowledge base. Architecturally distinct from any customer's identifiable data."

"What third-party AI providers do you use?"

"Enterprise-grade providers that are contractually prohibited from training on your data. We don't disclose specific providers because those relationships may evolve. What matters is the contractual protection — that's in our MSA."

"Do you build your own AI models from customer data?"

"Qumis develops specialized coverage analysis capabilities — classification, quality scoring, extraction — using de-identified, aggregated intelligence from our knowledge base. No



customer's identifiable data is used. This is how every serious vertical AI platform works — it's how we continuously improve accuracy while maintaining complete data privacy."

"What happens to our data if we leave?"

"All customer data is securely deleted within 60 days. Any de-identified intelligence derived from aggregated patterns during your time on the platform persists in the knowledge base — but by definition, that data doesn't identify you or contain your specific policies."

"Will you sign our AI data usage addendum?"

"Happy to review it. Our MSA Sections 8.5 and 4.4 already address data isolation, model training restrictions, third-party providers, and aggregated data. Let me share the relevant sections so your legal team can compare."

Language Discipline

Use	Avoid	Why
"Never used to train large language models"	"Never used to train AI" / "never used for training"	The commitment is specific to foundational LLMs. Broader language forecloses domain-specific capabilities built on de-identified data.
"Logically segregated"	"Physically isolated" / "dedicated environment"	Logical segregation is what you provide. Physical isolation implies infrastructure you don't offer.
"De-identified, aggregated intelligence"	"Anonymized data"	More precise. "Anonymized" has specific legal definitions that may not match your process.
"Proprietary knowledge base" / "coverage intelligence layer"	"Our database of customer policies"	Positions the asset correctly — as derived intelligence, not collected documents.



"Specialized domain capabilities"	"Our AI models trained on data"	Describes the same thing without triggering the training fear response.
"Processed through AI models"	"Fed into AI" / "sent to AI"	"Processed" is clinical. "Fed" sounds careless.
"Contractually prohibited"	"We promise" / "we don't do that"	Contractual language is stronger than verbal assurances.
"Enterprise-grade AI providers"	Naming specific providers	Relationships may change. Naming invites comparison to consumer products.
"Derived data structures"	"Graph database" / "knowledge graph" (customer-facing)	Technically accurate, doesn't attract scrutiny.

Red Lines: What Your Team Must Never Say

"We never use your data for anything except your own analysis." — Conflicts with MSA Section 4.4. If legal reads this statement and then reads the contract, trust collapses.

"We built our own large language model." — You use third-party foundation models with a proprietary intelligence layer. This claim invites scrutiny you can't survive.

"Your data never leaves our servers." — Customer data is processed through third-party providers for inference. This will be caught in any serious security review.

"We train our AI on customer data to make the platform better." — Conflates multiple distinct activities into the single statement buyers fear most. Indefensible in a sales context even though elements are technically true for the de-identified layer.

"Our AI is 100% accurate" / "Our AI never hallucinates." — Creates legal liability. Use citation-linked outputs and "I don't know" safeguards framing.



Internal Note

The "large language models" qualifier throughout this section is deliberate and must not be softened to "any AI" or "any models." The MSA prohibits training "general-purpose or foundational" large language models on customer data. It does not prohibit building specialized domain capabilities from de-identified aggregated data under Section 4.4. This precise language preserves full product optionality — graph databases, classifiers, evaluators, scoring models, SLMs — while maintaining an honest, defensible commitment to customers. If a customer's legal team pushes to narrow the aggregated data or permitted uses provisions, escalate to Dan.